

---

# **CSE 431**

## **Computer Architecture**

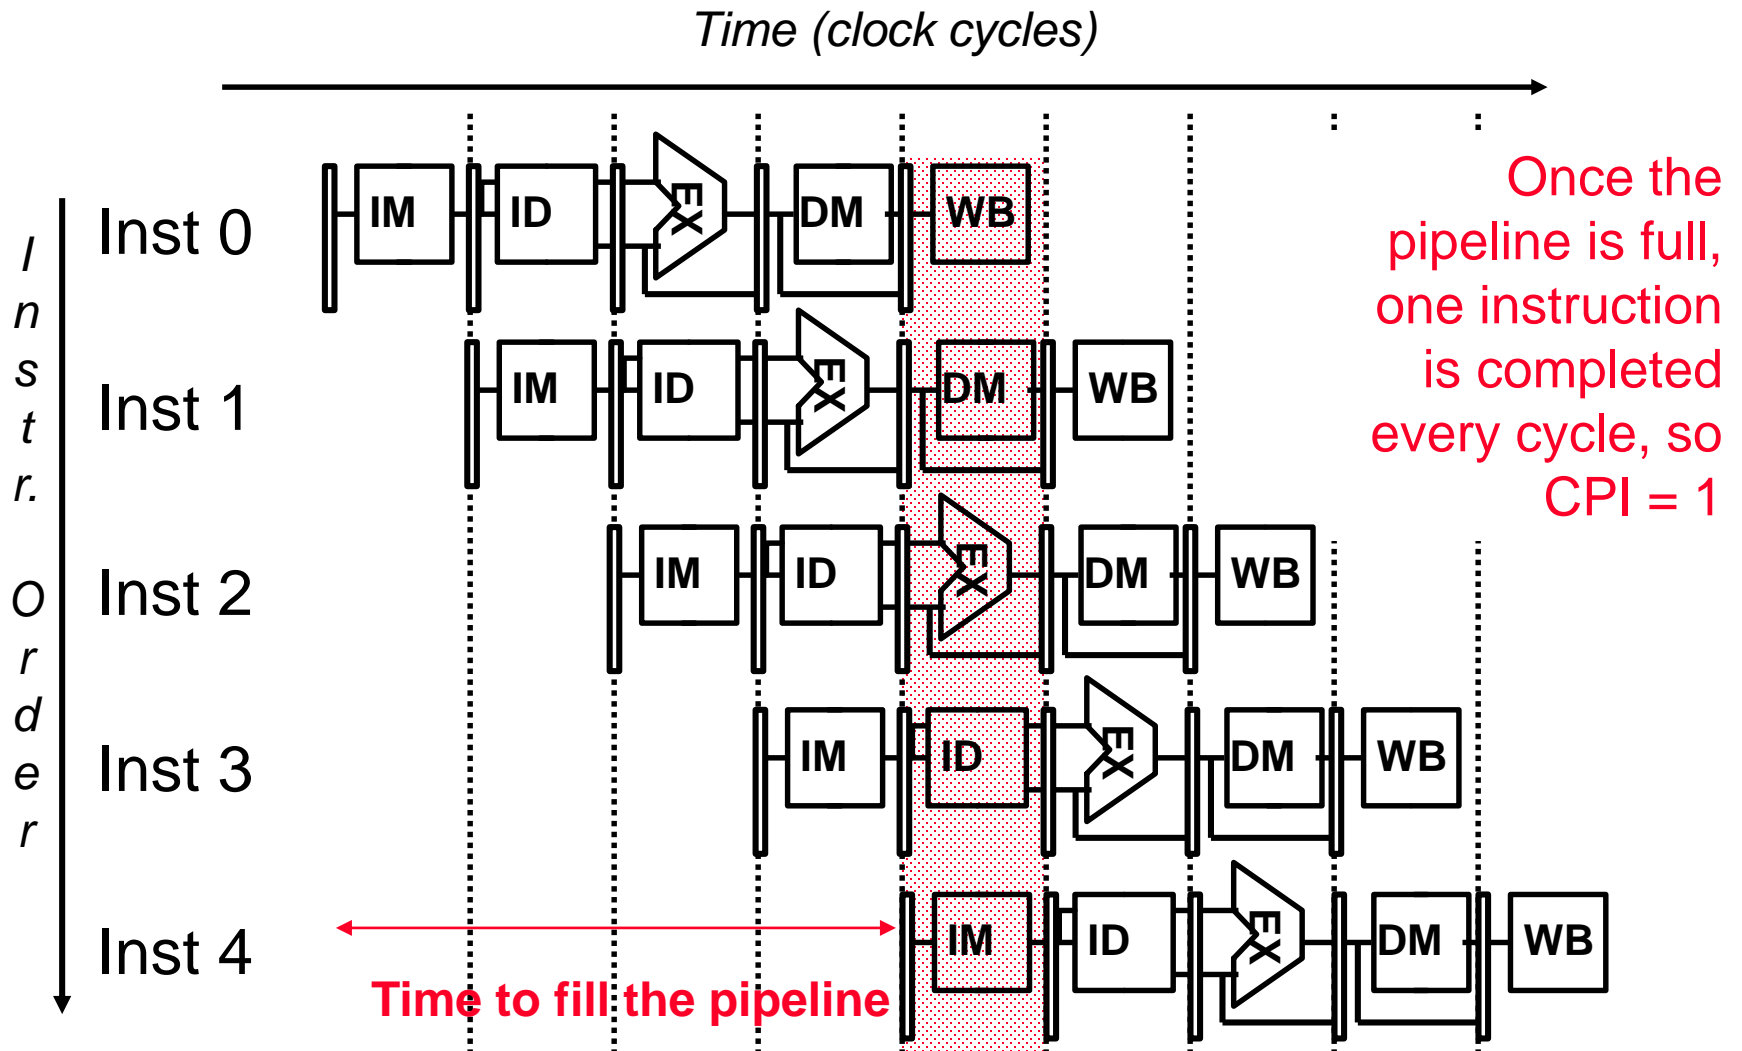
### **Fall 2017**

## **The Pipelined Processor: Dealing with Branches and Exceptions Part C**

Mahmut Taylan Kandemir ([www.cse.psu.edu/~kandemir](http://www.cse.psu.edu/~kandemir))

[Adapted from *Computer Organization and Design, 5<sup>th</sup> Edition*,  
Patterson & Hennessy, © 2014, Morgan Kaufmann]

# Review: Why Pipeline? For Performance!



# Review: Pipelining - What Makes it Hard ?

## ❑ Pipeline Hazards

- **structural hazards**: attempt to use the same resource by two different instructions at the same time
- **data hazards**: attempt to use data before it is ready
  - An instruction's source operand(s) are produced by a prior instruction still in the pipeline
- **control hazards**: attempt to make a decision about program control flow before the condition has been evaluated and the new PC target address calculated
  - branch and jump instructions, exceptions

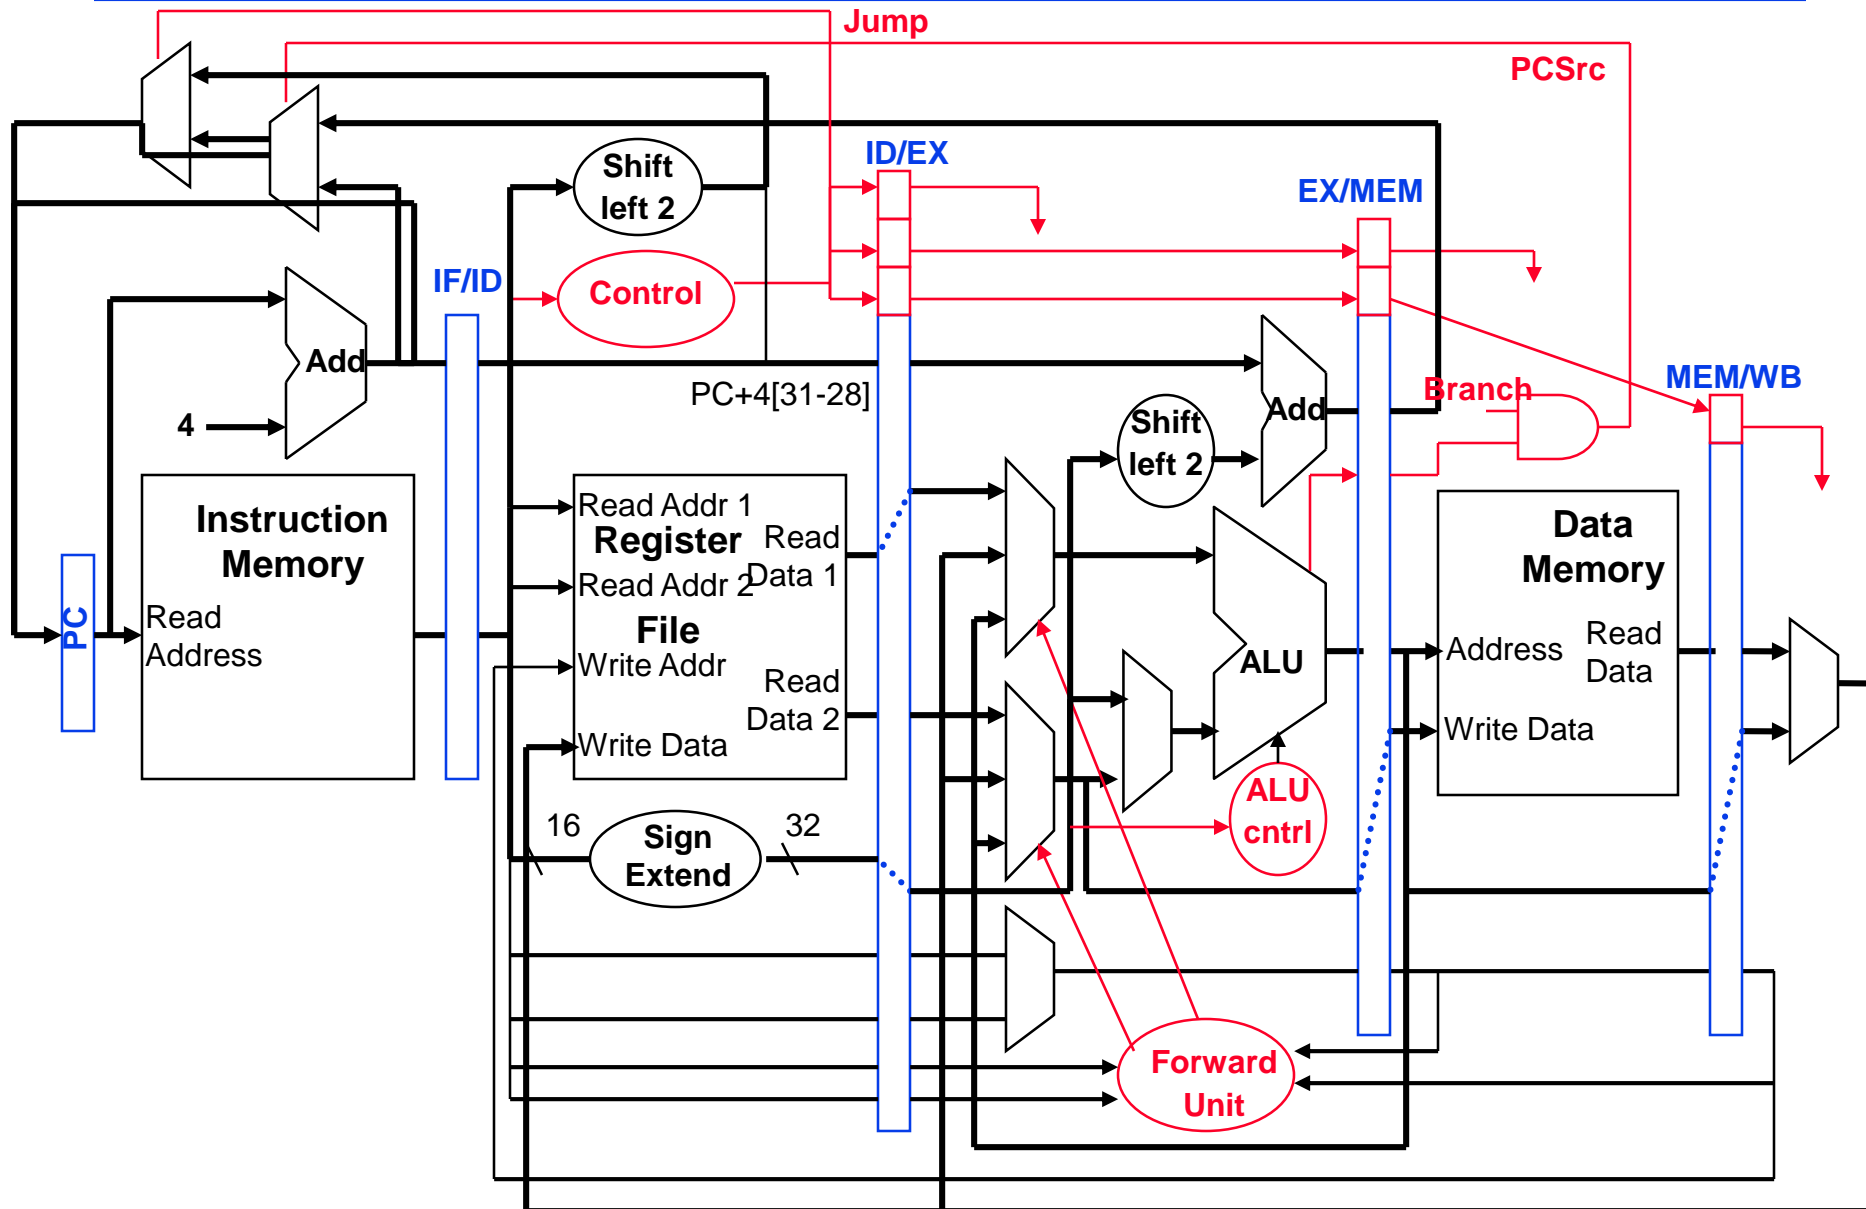
- ## ❑ Pipeline hardware control must **detect** the hazard and then take action to **resolve** hazard

# Control Hazards

---

- ❑ When the flow of instruction addresses is not sequential (i.e.,  $PC = PC + 4$ ); incurred by change of flow instructions
  - Unconditional branches (`j`, `jal`, `jr`)
  - Conditional branches (`beq`, `bne`)
  - Exceptions
- ❑ Possible approaches
  - Stall (impacts CPI)
  - Move decision point as early in the pipeline as possible, thereby reducing the number of stall cycles
  - Delay decision (requires compiler support)
  - Predict and hope for the best !
- ❑ Control hazards occur less frequently than data hazards, but there is nothing as effective against control hazards as forwarding is for data hazards

# Review: Datapath Branch and Jump Hardware



# Control (Branch, Jump) Hazards

## ❑ What do we need to know?

- Next instruction target address, maybe sequential (PC+4)  
*or*
  - `beq`, PC+4 + branch instruction's sign-extended offset which is computed during EX by the Shift Left 2 – Add logic
  - `j`, `jal`, constant address field read from IM during IF (26 bits)
  - `jr`, `jalr`, read from RF during ID (32 bits)
  - `trap` instruction or exception, obtained from table lookup in the OS (32 bits)
- Branch decision outcome (ALU zero flag)
  - continue sequentially? or jump to the branch target address?


## ❑ When do we need to know it?

- As early as possible in the pipeline

# Remember: Executing Branch Operations

## ❑ Branch operations involves

- compare the operands read from the RF during decode for equality (**zero** ALU output)
- compute the branch target address by adding the updated PC to the 16-bit signed-extended offset field in the instr



# Control (Branch, Jump) Hazards, con't

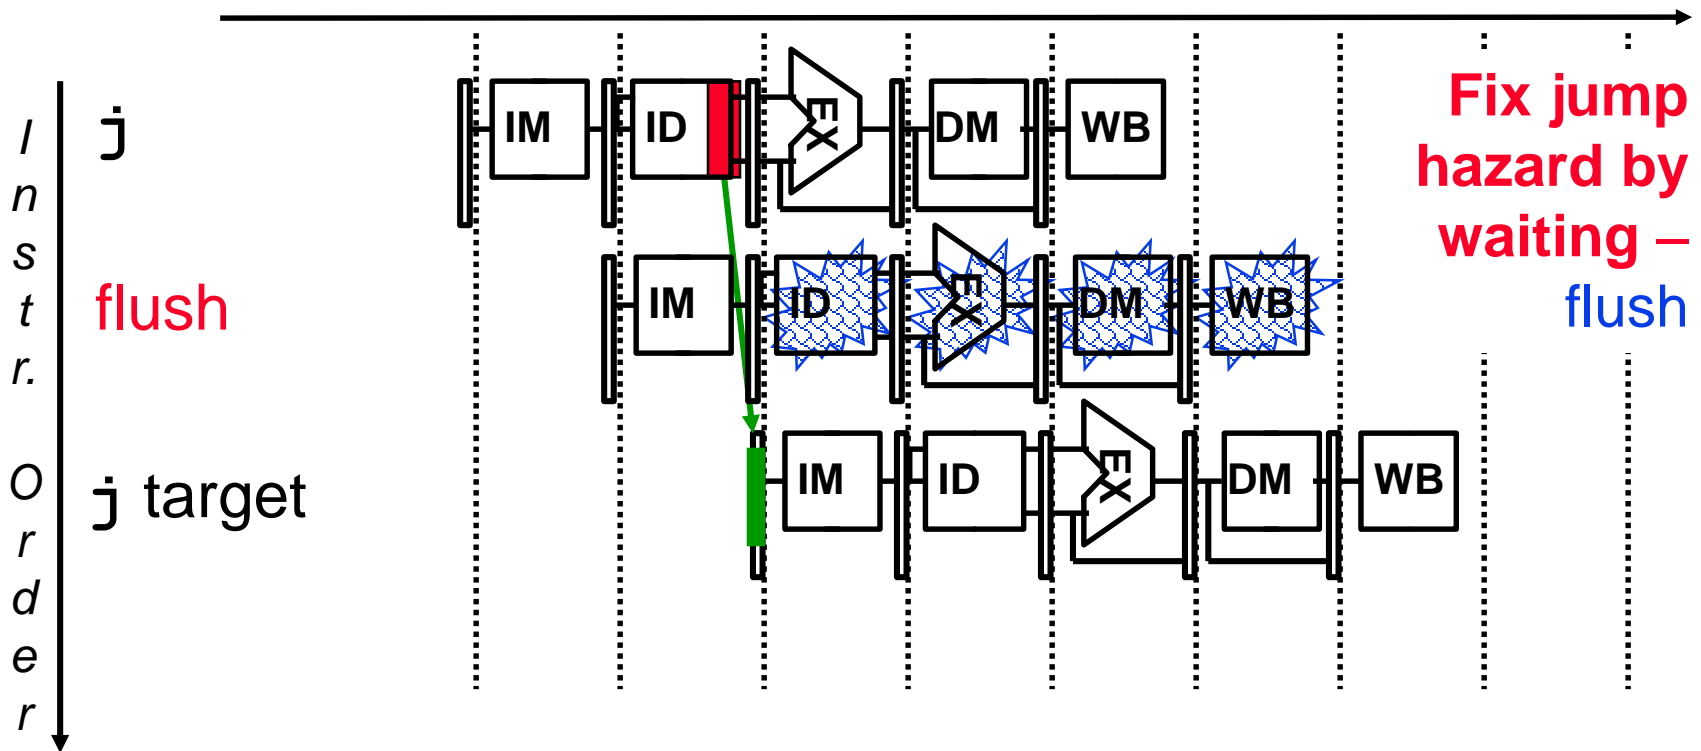
## ❑ When do we act on the decision?

- As soon as possible
  - we decided too late? We already fetched another instruction, and may need to discard it (flush it)
    - Maybe will have to flush more than one instruction?
  - we guessed, and were right – this is good
  - we guessed, and were wrong – now need to fix things
- Guesses require an evaluation of the success rate, by simulation (before the fact) or by measurement (after the fact)
  - can measurement improve the quality of the guesses?
  - recent history is often a reasonable predictor of the near future



# Jumps Incur One Stall

- ❑ Jumps not decoded until ID, so one **flush** is needed
  - To flush, set `IF/ID.Flush` to zero the bits in the IF/ID pipeline register (turning it into a `noop`)



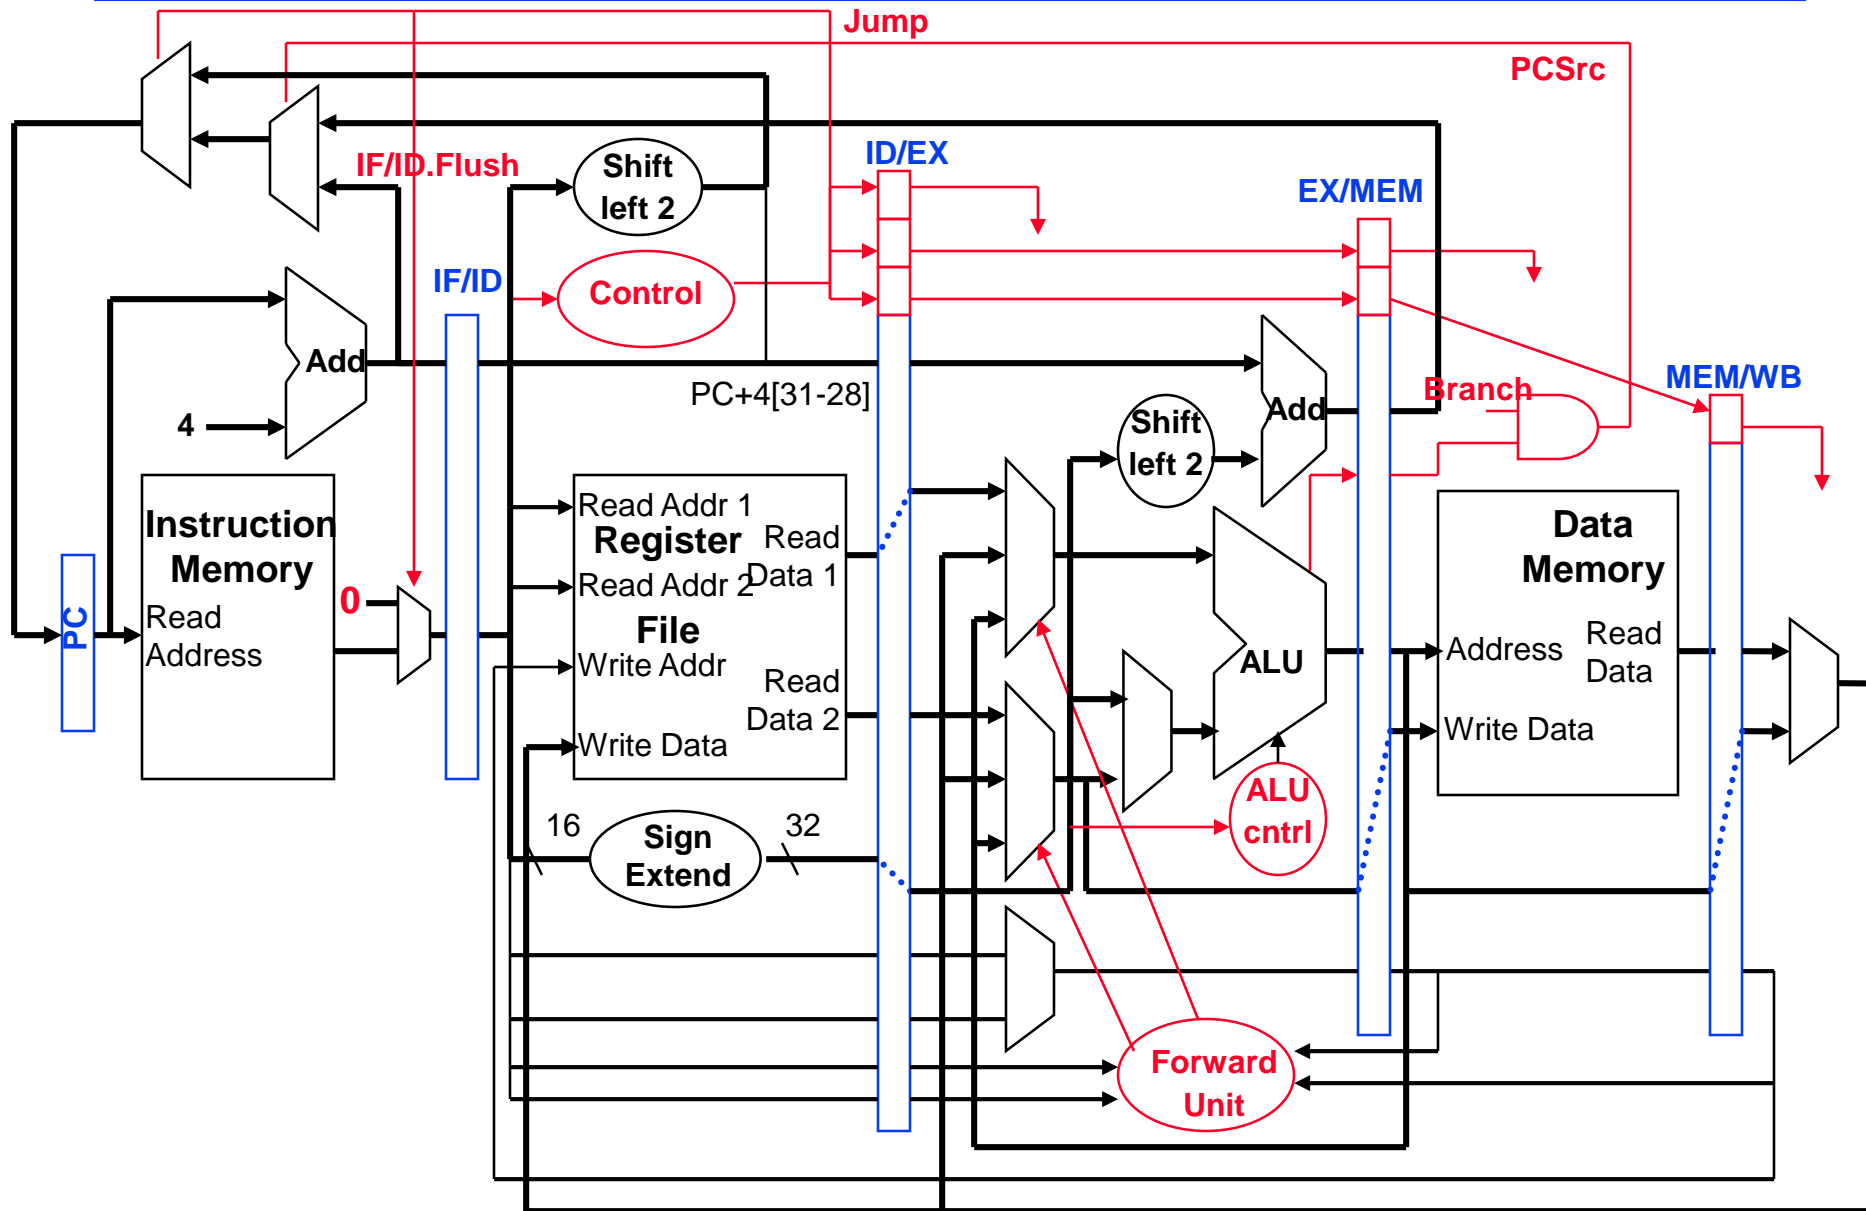
- ❑ Fortunately, jumps are very infrequent – only ~3% of the SPECint instruction mix

## Two “Types” of Stalls

---

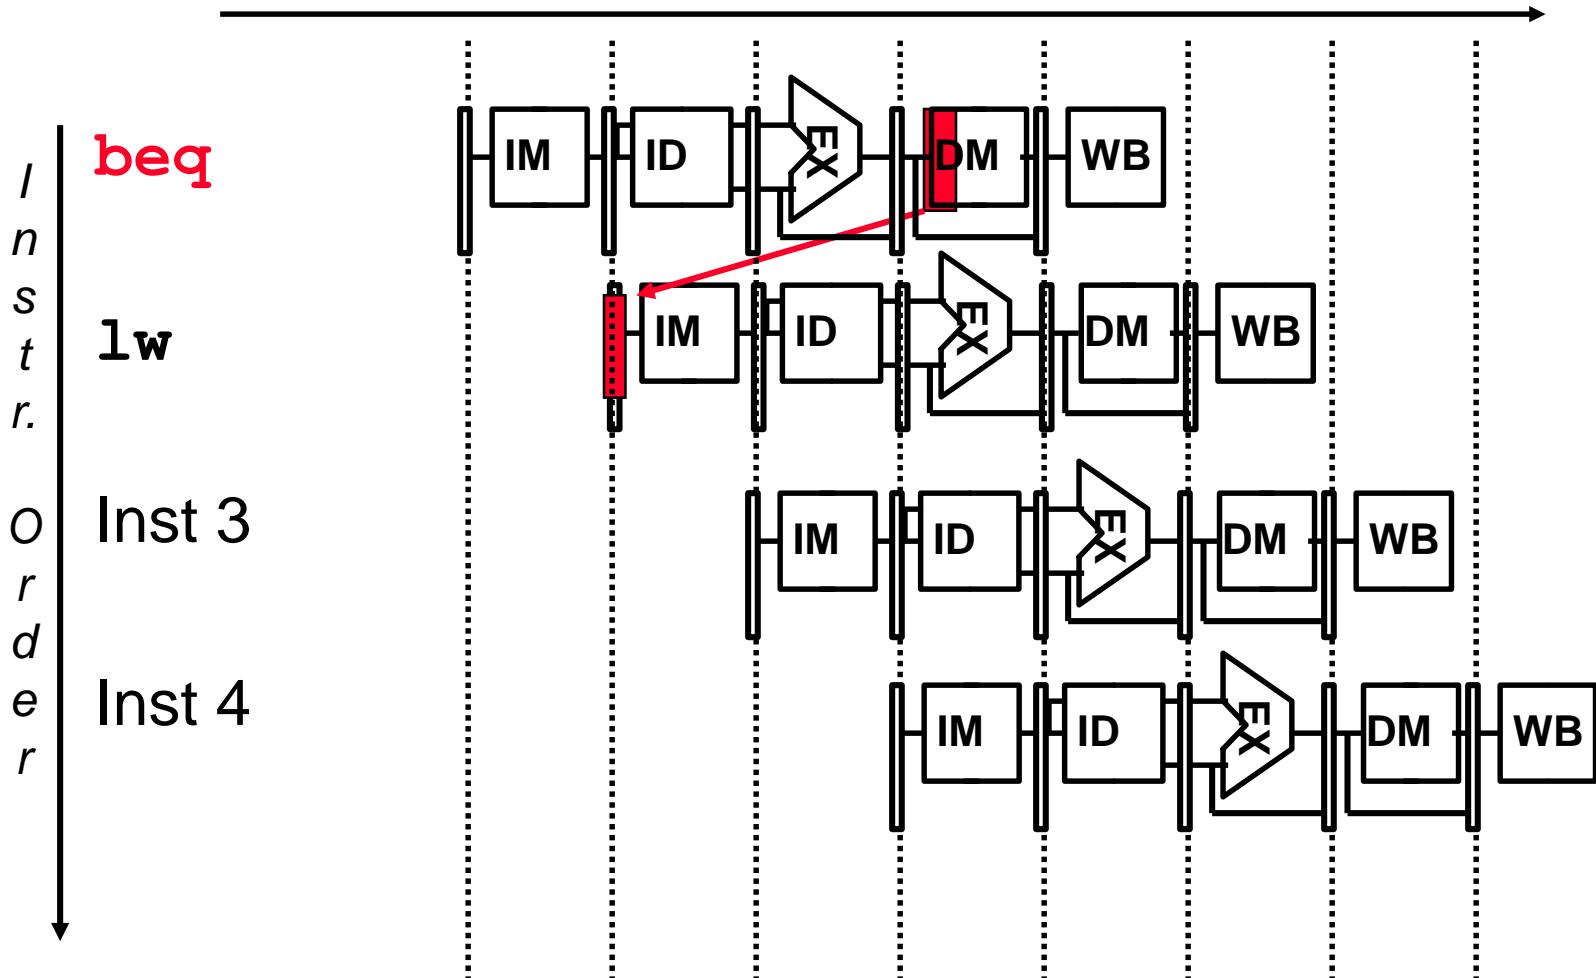
- ❑ **Stall** (or bubble) where a `noop` instruction is **inserted between** two instructions in the pipeline (as done for load-use)
  - Keep the instructions *earlier* in the pipeline (later in the code) from progressing down the pipeline for a cycle (“bounce” them in place with write control signals to the pipeline registers)
  - Insert `noop` by zeroing the control bits in the pipeline register at the appropriate stage (with `IF/ID.Bubble`)
  - Let the instructions later in the pipeline (earlier in the code) progress normally down the pipeline
- ❑ **Flush** (or instruction squashing) where an instruction in the pipeline is **replaced** with a `noop` instruction (as done for instructions located sequentially after `j` instructions with `IF/ID.Flush`)
  - Zero the control bits in the IF/ID pipeline register of the instruction to be flushed (the one just after the `j` instruction)

# Supporting ID Stage Jumps



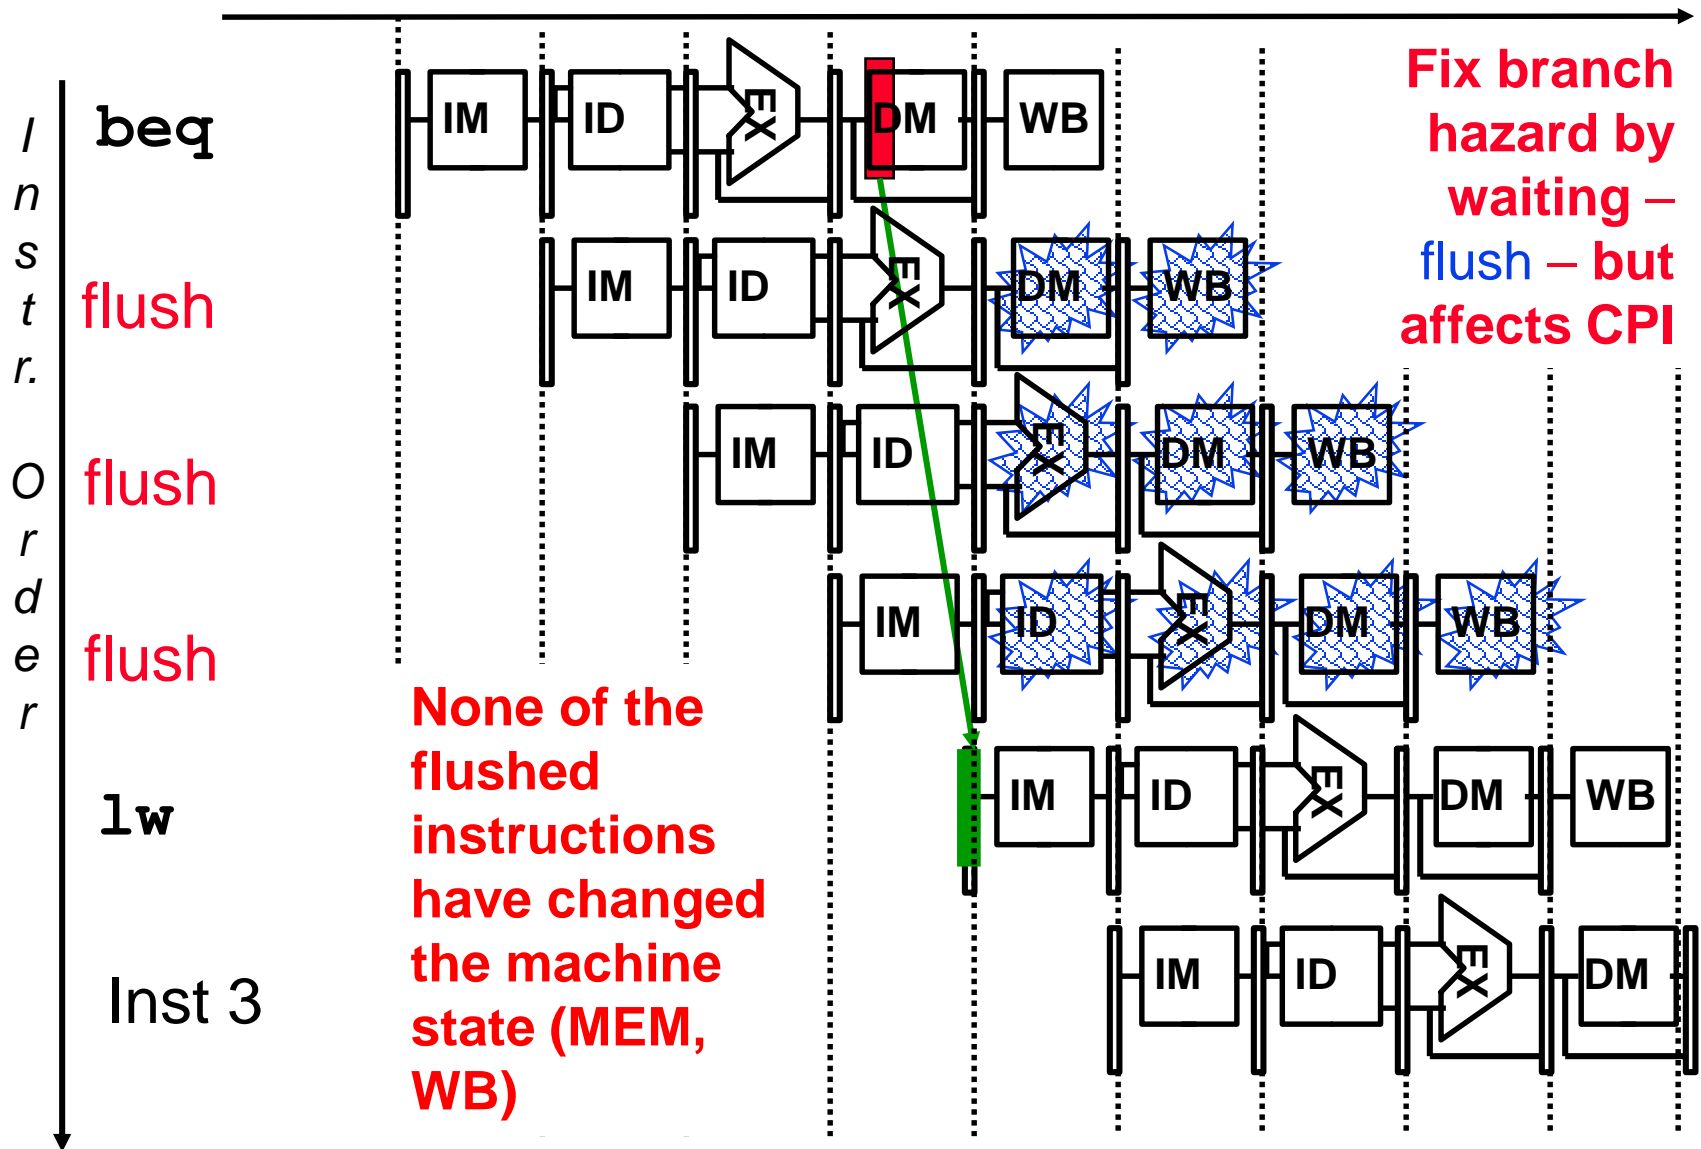
# Branch Instructions Cause Control Hazards

- Dependencies backward in time cause hazards



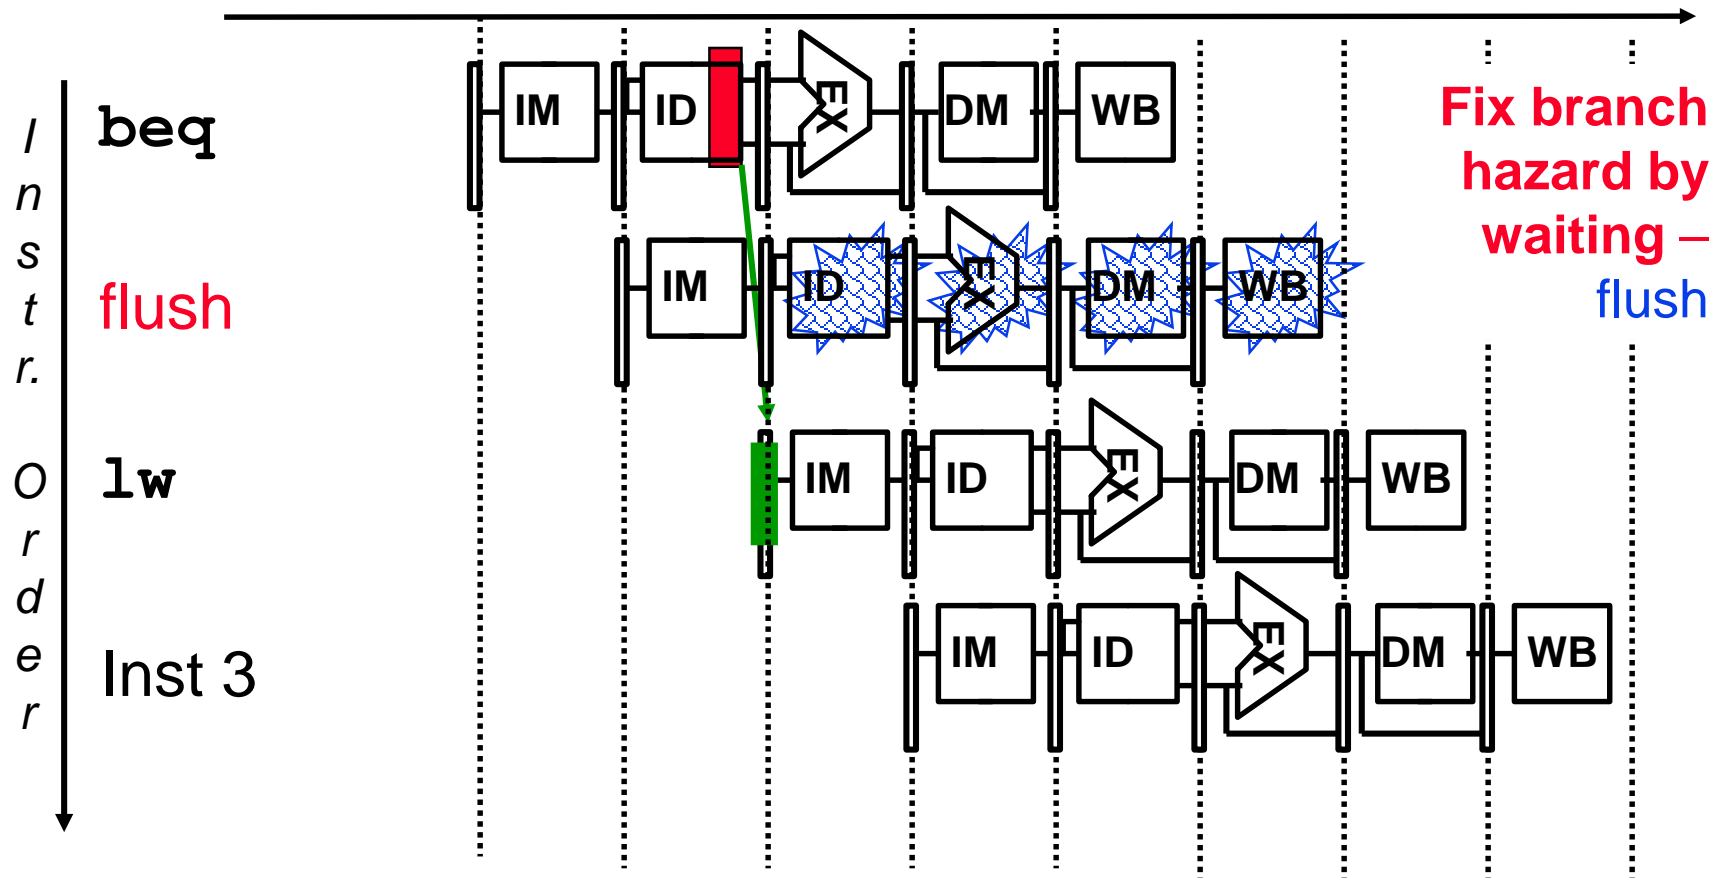
- Which instruction is executed after the branch ?

# One Way to “Fix” a Branch Control Hazard



# Another Way to “Fix” a Branch Control Hazard

- ❑ Move branch decision hardware back to as **early** in the pipeline as possible – i.e., during the decode cycle

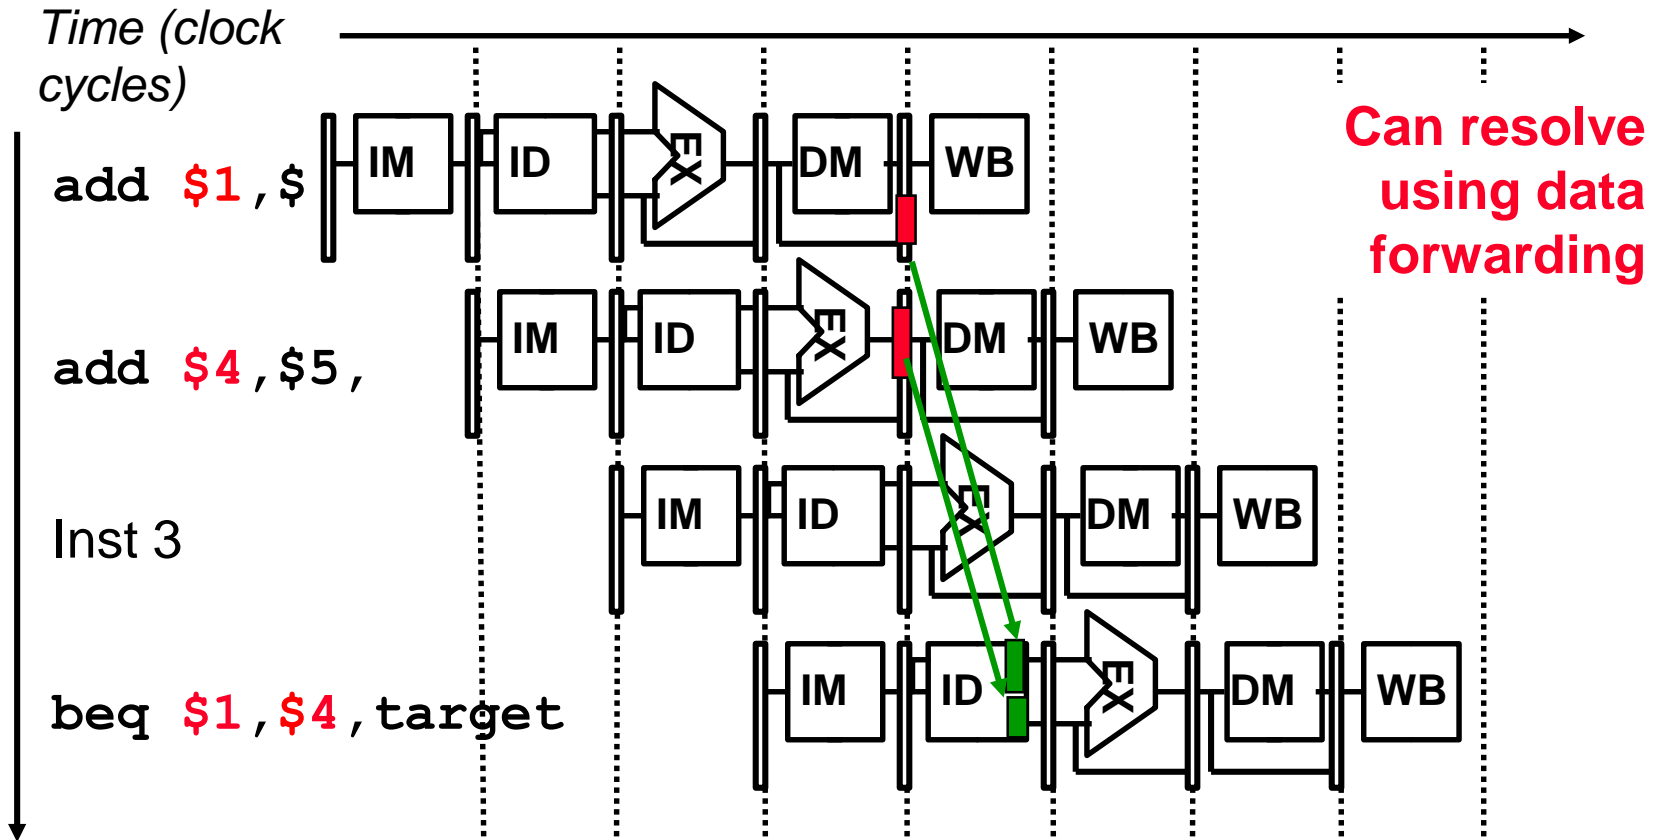


# Reducing the Delay of Branches

- ❑ Move the branch decision hardware back to the EX stage
  - Reduces the number of stall (flush) cycles to two
  - Adds an `and` gate and a `2x1 mux` to the EX timing path
- ❑ Move the branch decision hardware back to the ID stage
  - Reduces the number of stall (flush) cycles to one (like with jumps)
    - But now need to add **forwarding hardware** in ID stage
  - Computing branch target address can be done in parallel with RegFile read (done for all instructions – only used when needed)
  - Comparing the registers can't be done until after RegFile read, so adds a comparator to the ID timing path
- ❑ For deeper pipelines, branch decision points can be even later in the pipeline, incurring more stalls

# Yet Another Complication

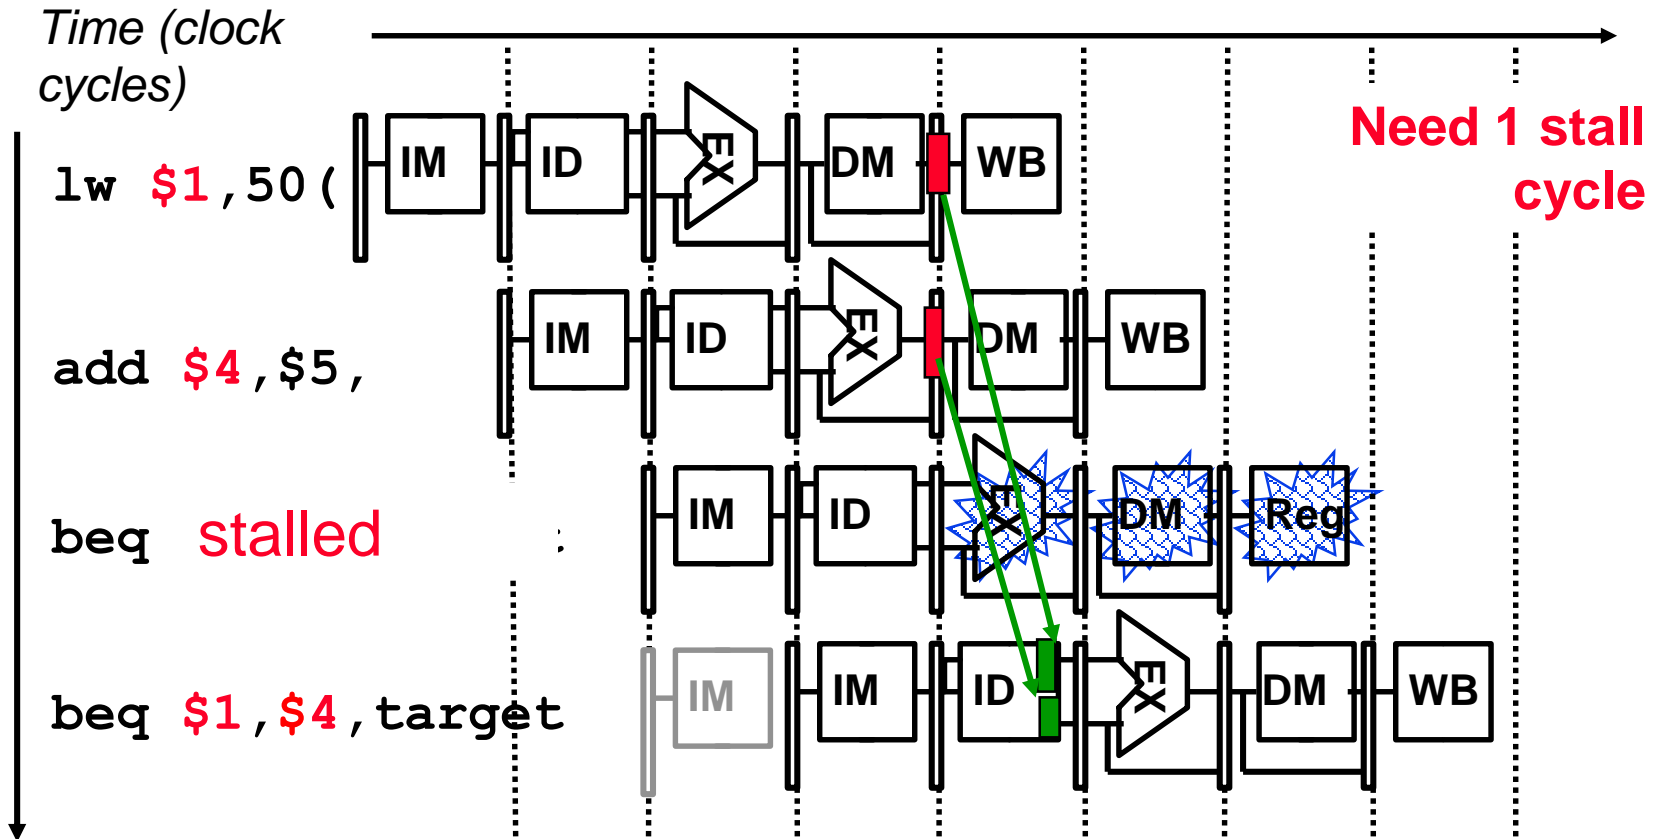
- What if a comparison register is a destination of the 2<sup>nd</sup> or 3<sup>rd</sup> preceding ALU instruction ?





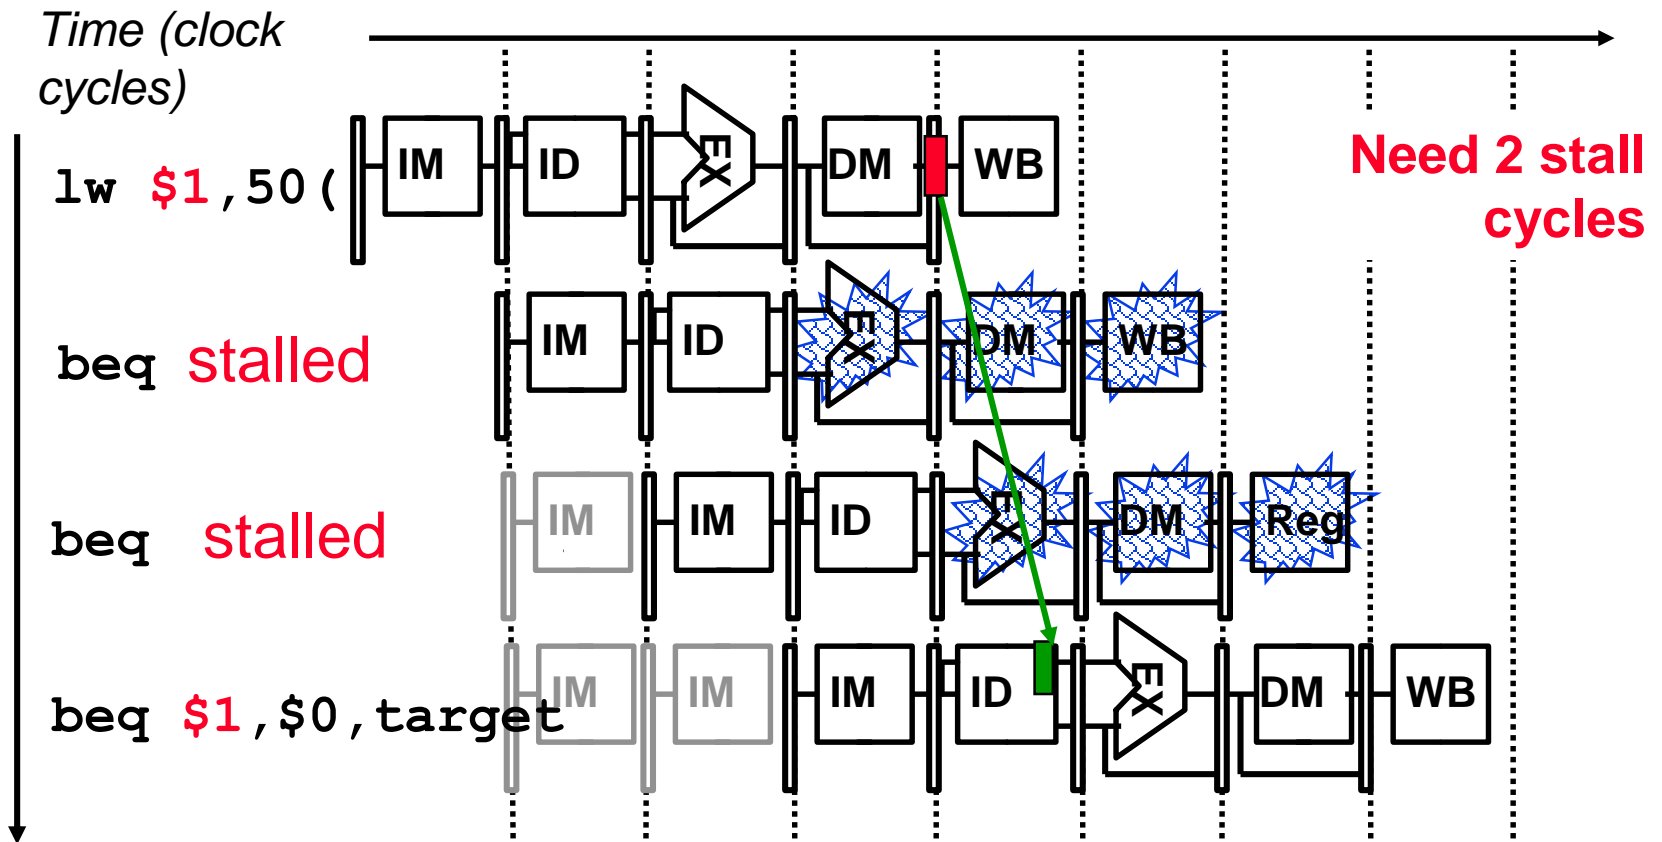
# And Another Branch Data Hazard

- What if a comparison register is a destination of the preceding ALU instruction or the 2<sup>nd</sup> preceding  $lw$  instr?



# And Yet Another Branch Data Hazard

- What if a comparison register is a destination of the immediately preceding preceeding  $lw$  instr?



# Static Branch Prediction

---

- ❑ Resolve branch hazards by **assuming** a given outcome and proceeding without waiting to see the actual branch outcome
- 1. **Predict not taken** – *always* predict branches will **not** be taken, continue to fetch from the sequential instruction stream, only when branch *is* taken does the pipeline stall
  - If taken, **flush** instructions **after** the branch (earlier in the pipeline, later in the code)
    - in IF, ID, and EX stages if branch logic in MEM – **three** stalls
    - In IF and ID stages if branch logic in EX – **two** stalls
    - in IF stage if branch logic in ID – **one** stall
  - ensure that those flushed instructions haven't changed the machine state – automatic in the MIPS pipeline since machine state changing operations are at the tail end of the pipeline (MemWrite (in MEM) or RegWrite (in WB))
  - restart the pipeline at the branch destination

# Static Branching Structures

---

- ❑ Always predict NT works well for “top of the loop” branching structures

- But such loops have jumps at the bottom of the loop to return to the top of the loop – and incur the jump stall overhead every time through the loop


```
Loop: beq $1,$2,Out
      1st loop instr
      .
      .
      .
      last loop instr
      j Loop
Out:  fall out instr
```

- ❑ Always predict NT doesn't work well for “bottom of the loop” branching structures

- Guess wrong every time through the loop except the last time (when we fall out of the loop)

```
Loop: 1st loop instr
      2nd loop instr
      .
      .
      .
      last loop instr
      bne $1,$2,Loop
      fall out instr
```

# 0-Bit Predictors, Loop for 10 Iterations



iteration	predict	actual	predict	actual
1	NT	T	T	T
2	NT	T	T	T
3	NT	T	T	T
4	NT	T	T	T
5	NT	T	T	T
6	NT	T	T	T
7	NT	T	T	T
8	NT	T	T	T
9	NT	T	T	T
10	NT	NT	T	NT
10% accuracy		90% accuracy		

```
pre-loop instr
Loop:
  1st loop instr
  2nd loop instr
    .
    .
    .
  last loop instr
  bne $1,$2,Loop
post-loop instr
```

So use  
predict  
Taken  
(T)?


# Other Branch Prediction Possibilities

---

2. **Predict taken (T)** – predict branches will always be taken
  - BUT predict taken *always* incurs one stall cycle (*if* branch destination hardware has been moved to the ID stage)
  - Is there a way to “cache” the **address** of the branch target instruction, or *better yet* the **actual** branch target **instruction** itself ?? Yes ... stay tuned !
- ❑ As the branch penalty increases (for deeper pipelines), a simple static prediction scheme will hurt performance. With more hardware, it is possible to try to predict branch behavior **dynamically** during program execution
3. **Dynamic branch prediction** – predict branches at run-time using *run-time* information

# Dynamic Branch Prediction Buffer

- ❑ A **branch prediction buffer** (aka branch history table (**BHT**)) in the IF stage addressed by the low order bits of the PC, contains bit(s) (passed to the ID stage through the IF/ID pipeline register) that tell whether the branch was taken or not the last time it was execute



# Branch History Table

---

- ❑ The BHTs prediction bits may predict incorrectly (may be a wrong prediction for this branch this iteration or may be from a different branch with the same low order PC bits) but the doesn't affect **correctness**, just **performance**
  - Branch decision occurs in the ID stage after determining that the fetched instruction is a branch and checking the prediction bit(s)
- ❑ If the prediction is wrong, flush the incorrect instruction(s) in pipeline, restart the pipeline with the right instruction, and invert the prediction bit(s)
- ❑ Default BHT size in SimpleScalar is 2048 entries (11 low order bits of the PC)
  - A 4096-bit BHT using 2-level adaptive prediction varies from 1% misprediction (nasa7, tomcatv) to 18% (eqntott) misprediction rate



# 1-bit NT Dynamic Branch Prediction Accuracy

## ❑ A 1-bit NT dynamic predictor will be incorrect twice

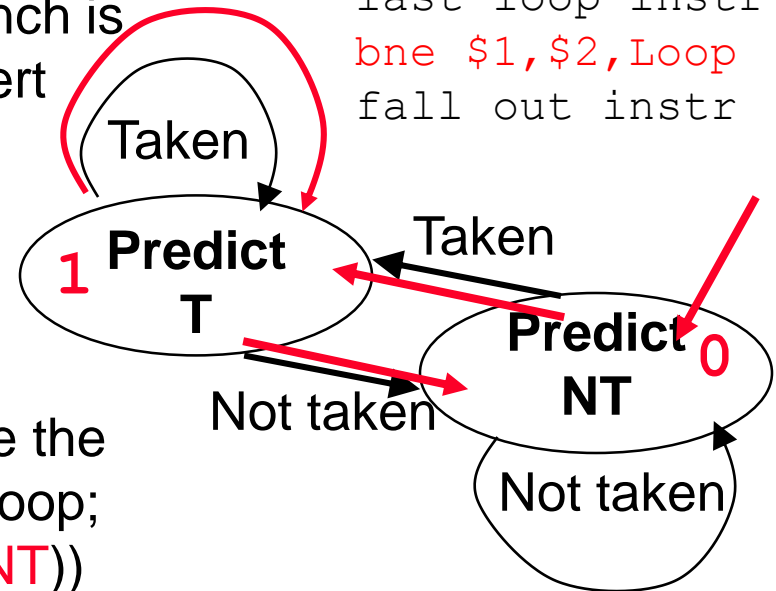
- Assume predict\_bit = 0 (NT) to start (indicating branch not taken) and loop control is at the bottom of the loop code

Loop: 1<sup>st</sup> loop instr  
2<sup>nd</sup> loop instr

⋮


last loop instr  
bne \$1,\$2,Loop  
fall out instr

1. First time through the loop, the predictor mispredicts the branch since the branch is taken back to the top of the loop; invert prediction bit (predict\_bit = 1 (T))
2. As long as branch is taken (looping), prediction is correct
3. Exiting the loop, the predictor again mispredicts the branch since this time the branch is not taken falling out of the loop; invert prediction bit (predict\_bit = 0 (NT))



- ## ❑ For 10 times through the loop we have a 80% prediction accuracy for a branch that is taken 90% of the time

# 1-Bit Dynamic Predictor, Loop for 10 Iterations

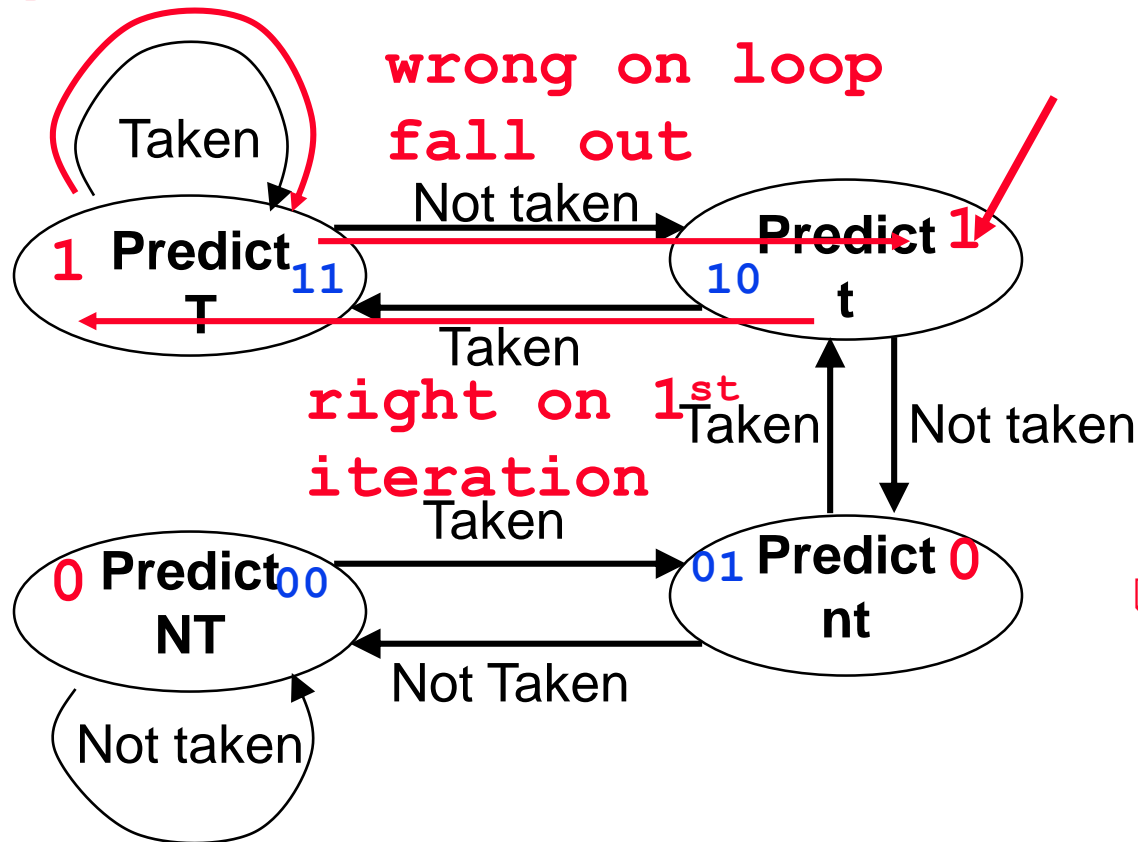


	Initial State = 0			Initial State = 1		
iteration	state	predict	actual	state	predict	actual
1	0	NT	T	1	T	T
2	1	T	T	1	T	T
3	1	T	T	1	T	T
4	1	T	T	1	T	T
5	1	T	T	1	T	T
6	1	T	T	1	T	T
7	1	T	T	1	T	T
8	1	T	T	1	T	T
9	1	T	T	1	T	T
10	1	T	NT	1	T	NT
	0	80% accuracy		0	90% accuracy	

## 2-bit Dynamic Branch Predictors

- ❑ A 2-bit scheme can give 90% accuracy since a prediction must be wrong twice before the prediction bit is changed

right 9 times

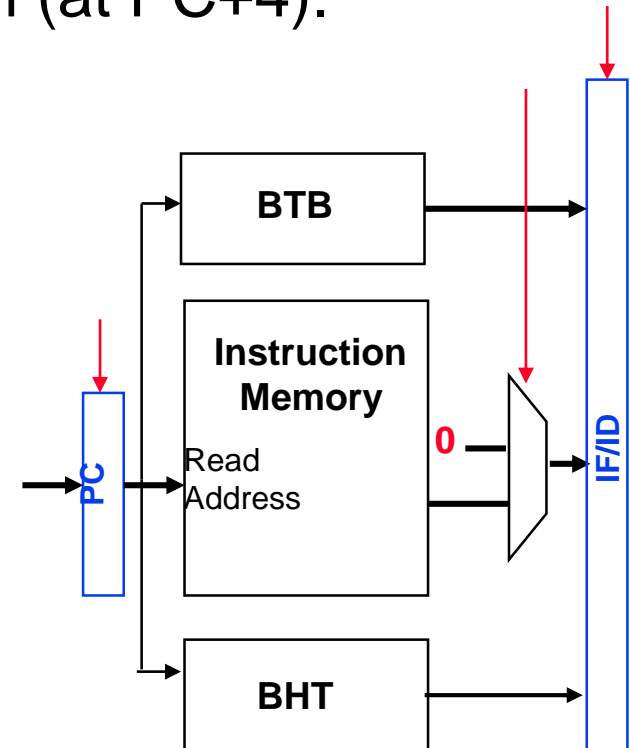


```
Loop: 1st loop instr
      2nd loop instr
      .
      .
      .
      last loop instr
      bne $1,$2,Loop
      fall out instr
```

- ❑ The BHT also stores the initial FSM state

# But Also Must Know the Branch Target

- ❑ A **branch target buffer (BTB)** in the IF stage can cache the branch target **address**, but remember we also need to fetch the next sequential instruction (at PC+4).
  - Would need a two read port IM
- ❑ Or, the BTB can cache the actual branch target **instruction** while the Instruction Memory is fetching the next sequential instruction
  - Would need a **two** read port IM
  - ID stage can then select between PC+4 and branch target instruction
- ❑ If we predict correctly, stalls can be avoided no matter which direction the branch goes



# Summary of 2-bit Dynamic Branch Predictors

---

- ❑ A 2-bit dynamic branch prediction scheme can give 90% accuracy since a prediction must be wrong twice before the prediction is changed
  - In a counter implementation, the counters are incremented when a branch is taken and decremented when not taken (and **saturate** at 00 or 11).
- ❑ BHT stores the initial state of the predictor's Finite State Machine (usually the last state last time through the loop)
- ❑ BTB stores the branch target instruction which is “fetched” along with the sequential instruction in the Fetch stage (when the branch is in the Decode stage)
- ❑ Since we read the prediction bits on every cycle, a 2-bit predictor will need both a read and a write access port for updating the prediction bits.

# State-of-the-Art in Branch Prediction

---

## ❑ Basic 2-bit predictor:

- For each branch:
  - Predict taken or not taken
  - If the prediction is wrong two consecutive times, change prediction

## ❑ Correlating predictor:

- Multiple 2-bit predictors for each branch
- One for each possible combination of outcomes of preceding  $n$  branches


## ❑ Local predictor:

- Multiple 2-bit predictors for each branch
- One for each possible combination of outcomes for the last  $n$  occurrences of this branch

## ❑ Tournament predictor:

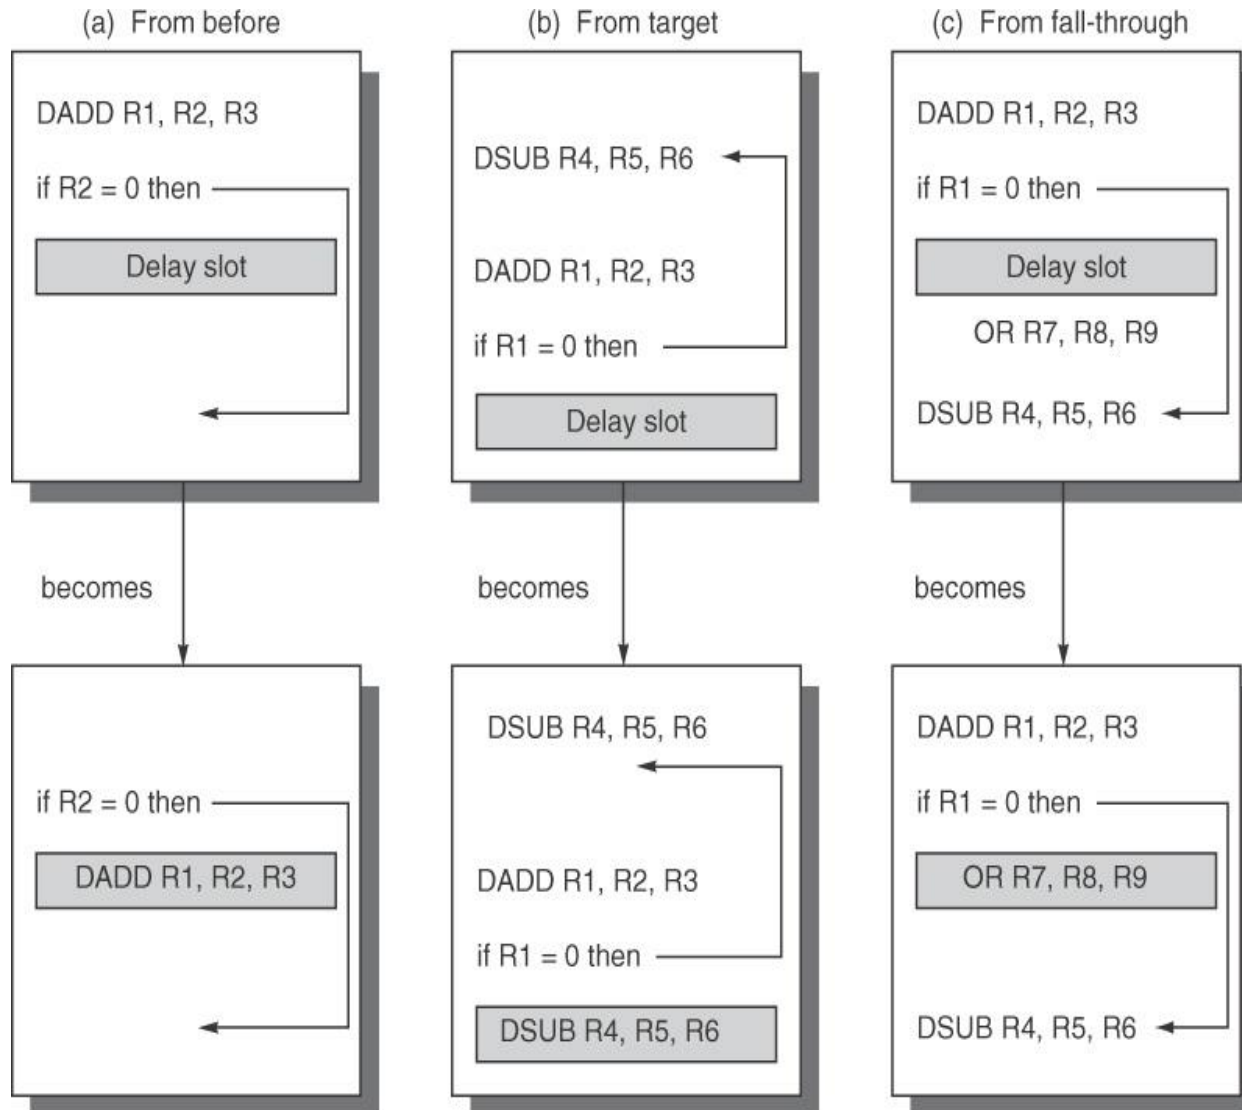
- Combine correlating predictor with local predictor

# Branch Prediction Performance



Branch predictor performance

# Branch Delay Slots





# Example

---

Consider a branch that is taken 80% of the time. On average, how many stalls are introduced for this branch for each approach below:

- Stall fetch until branch outcome is known
- Assume not-taken and squash if the branch is taken
- Assume a branch delay slot
  - You can't find anything to put in the delay slot
  - An instr before the branch is put in the delay slot
  - An instr from the taken side is put in the delay slot
  - An instr from the not-taken side is put in the slot

# Summary

---

- Techniques to handle control hazard stalls:
  - for every branch, introduce a stall cycle (note: every 6<sup>th</sup> instruction is a branch on average!)
  - assume the branch is not taken and start fetching the next instruction – if the branch is taken, need hardware to cancel the effect of the wrong-path instructions
  - predict the next PC and fetch that instr – if the prediction is wrong, cancel the effect of the wrong-path instructions
  - fetch the next instruction (branch delay slot) and execute it anyway – if the instruction turns out to be on the correct path, useful work was done – if the instruction turns out to be on the wrong path, hopefully program state is not lost

# SimpleScalar Branch Prediction

---

-bpred<type>	
not taken	Always predict not taken
taken	Always predict taken
perfect	Perfect prediction (but can't build it)
bimod	Bimodal predictor using a 2048 entry BHT with 2-bit counters and a BTB (512 sets, 4-way associativity)
2lev	2-level adaptive predictor
comb	Combined predictor (bimodal and 2-level adaptive)

- ❑ Many more possibilities, see [http://en.wikipedia.org/wiki/Branch\\_predictor](http://en.wikipedia.org/wiki/Branch_predictor)
- ❑ Prediction accuracy improves as the branch predictor grows in complexity

# Dealing with Exceptions

---


- ❑ Exceptions (aka interrupts) are just another form of control hazard. Exceptions arise from
  - R-type arithmetic overflow
  - Trying to execute an undefined instruction
  - An I/O device request
  - An OS service request (e.g., a page fault, TLB exception)
  - A hardware malfunction
- ❑ The pipeline has to stop executing the offending instruction in midstream, let all prior instructions complete, flush all following instructions, set a register to show the cause of the exception, save the address of the offending instruction, and then jump to a prearranged address (the address of the exception handler code)
- ❑ The software (OS) looks at the cause of the exception and “deals” with it

# Two Types of Exceptions

---

- ❑ Interrupts – asynchronous to program execution
  - caused by **external events**
  - may be handled **between** instructions, so can let the instructions currently active in the pipeline *complete* before passing control to the OS interrupt handler
  - simply suspend and resume user program
  
- ❑ Traps (Exception) – synchronous to program execution
  - caused by **internal events**
  - condition must be remedied by the trap handler for **that** instruction, so must stop the offending instruction *midstream* in the pipeline and pass control to the OS trap handler
  - the offending instruction may be retried (or simulated by the OS) and the program may continue or it may be aborted

# Where in the Pipeline Exceptions Occur




	Stage(s)?	Synchronous?
<input type="checkbox"/> Arithmetic overflow	EX	yes
<input type="checkbox"/> Undefined instruction	ID	yes
<input type="checkbox"/> TLB or page fault	IF, MEM	yes
<input type="checkbox"/> I/O service request	any	no
<input type="checkbox"/> Hardware malfunction	any	no

**Be aware that multiple exceptions can occur simultaneously in a *single* clock cycle!**

# Additions to MIPS to Handle Exceptions

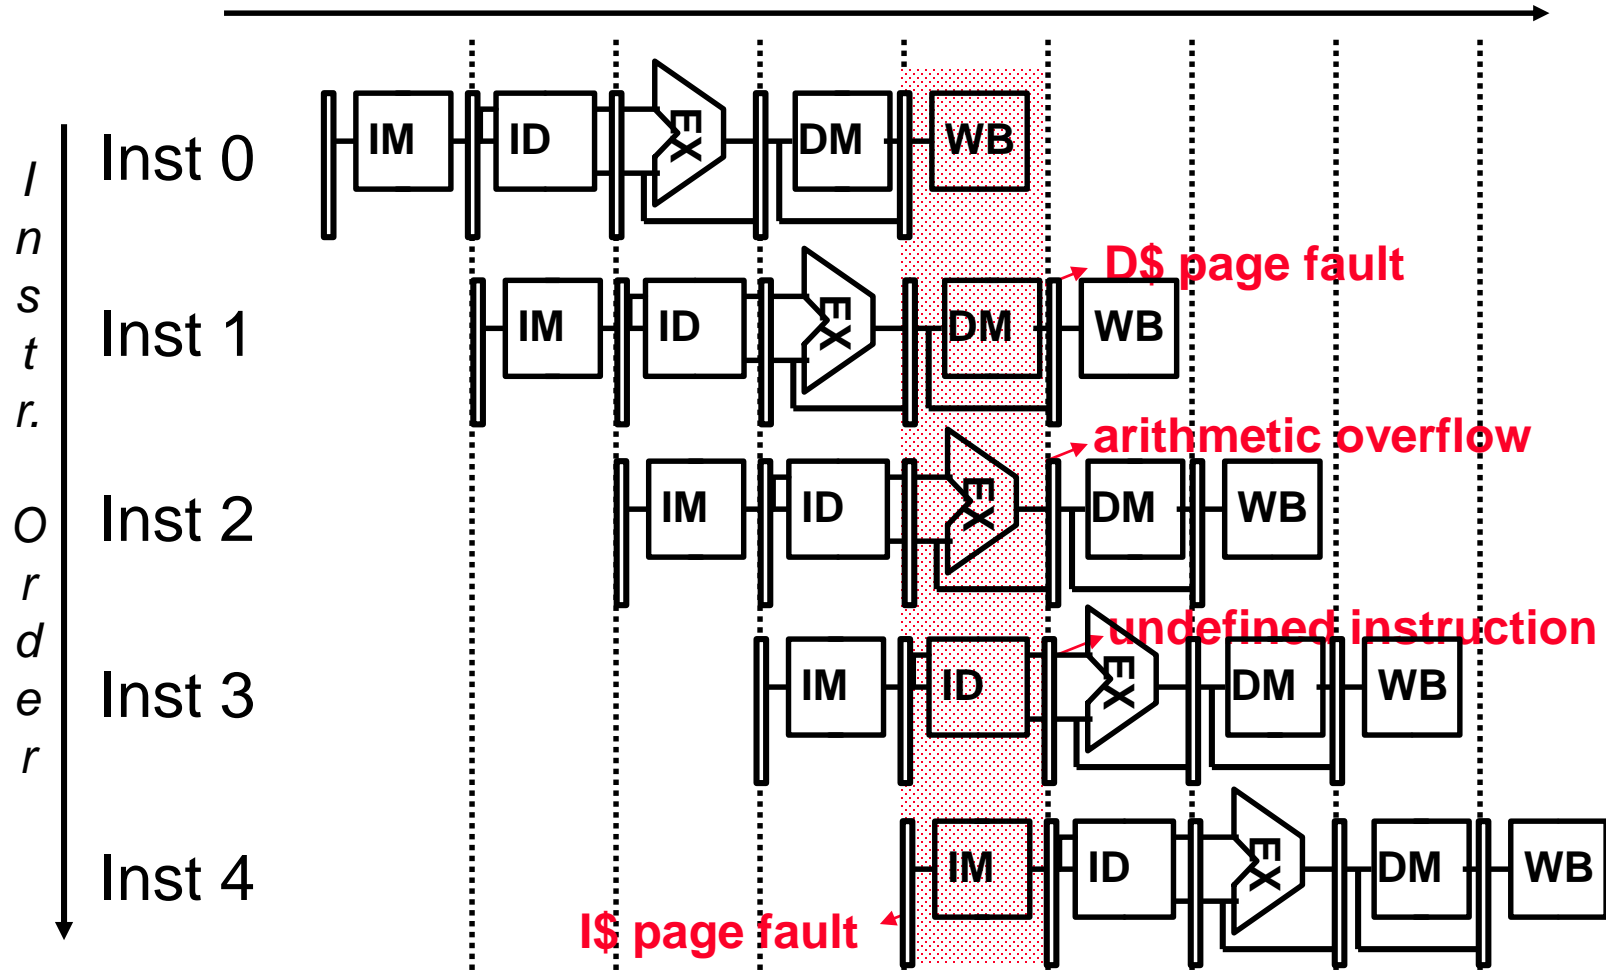
- ❑ Cause register (records exceptions) – hardware to record in Cause the exceptions and a signal to control writes to it (CauseWrite)
- ❑ EPC register (records the addresses of the offending instructions) – hardware to record in EPC the address of the offending instruction and a signal to control writes to it (EPCWrite)
  - Exception software must match exception to instruction
- ❑ A way to load the PC with the address of the exception handler
  - Expand the PC input mux where the new input is hardwired to the exception handler address - (e.g., 8000 0180<sub>hex</sub> for arithmetic overflow, 8000 0000<sub>hex</sub> for undefined instruction)
- ❑ A way to flush offending instruction and the ones that follow it

# Pipeline with Exception Extensions





# Multiple Simultaneous Exceptions



- ❑ Hardware sorts the exceptions so that the earliest instruction (D\$ page fault) is the one interrupted first

# Summary

---

- ❑ All modern day processors use pipelining for performance (a CPI of 1 and fast a CC)
- ❑ Pipeline clock rate limited by **slowest** pipeline stage – so designing a balanced pipeline is important
- ❑ Must detect and resolve hazards
  - Structural hazards – resolved by designing the pipeline correctly
  - Data hazards
    - Stall (impacts CPI)
    - Forward (requires hardware support)
  - Control hazards – put the branch decision hardware in as early a stage in the pipeline as possible
    - Stall (impacts CPI)
    - Delay decision (requires compiler support)
    - Static and **dynamic prediction** (requires hardware support)
- ❑ Pipelining complicates exception handling

# Reading Assignment

---

- ❑ Read Sections 4.8 and 4.9 from HP